

From Adaptive K to Routed RTL: A CHIA Loop for HEVC RDO Co-Design

Nguyen Hieu Duong · Duy Hieu Bui · Vietnam National University · 24020497@vnu.edu.vn

Artifact: <https://github.com/duongnguyenhieu/CHIA-RDO>, release v1.0.0

Abstract

HEVC intra coding exposes 35 luma prediction modes, making full rate-distortion optimization (RDO) expensive. Yet candidate count alone is not a hardware objective: an FPGA with P lanes executes $\lceil K/P \rceil$ batches, while replication changes area, routing, timing, and power. We propose CHIA-RDO, an evidence-preserving CHIA loop that joins pinned HM-16.20 experiments, a history-aware proposal controller, analytical screening, bit-exact Verilator regression, Pareto updates, GCP execution, and Vivado implementation. A frozen Adaptive- K policy reduces held-out RDO evaluations by **38.21%** while retaining **97.91%** of Full-RDO winners; its aggregate Y-PSNR BD-rate is -0.026% , interpreted as no observed penalty on this limited corpus. Physical feedback first exposed a 170-level combinatorial baseline and then an 80-level serial winner scan. Resource sharing, a depth-3 quant/dequant split, corpus-qualified 31-bit cost, and a balanced reduction tree yield a routed 16-lane kernel at a derived **68.120 MHz** and **8.134 M candidates/s**, $2.128\times$ the P1 baseline proxy. The design does *not* close 200 MHz, and kernel rate is not encoder throughput.

1. Problem and Opportunity

Stock HM ranks all 35 modes by rough mode decision (RMD), keeps a PU-size dependent shortlist, inserts most-probable modes, and performs expensive RDO. Prior work shows that content-adaptive pruning can reduce this work [2]. The cross-layer problem is that two policies with equal mean K may have different hardware cost because latency is quantized by $\lceil K/P \rceil$. Conversely, adding lanes may reduce batches but fail placement or deepen a routing-critical cone.

CHIA-RDO therefore optimizes a contract spanning three evidence classes: software coding quality, cycle-accurate RTL correctness, and routed physical results. Analytical estimates are retained but never relabeled as measured hardware.

ware.

Contributions

1. A composable CHIA graph connecting policy search, matched HM curves, RTL checks, cloud workers, persistent history, and backend feedback.
2. A frozen, independently specified Adaptive- K policy, evaluated separately from later hardware-aware policies.
3. A physically driven RTL evolution from monolithic replication to serial shared lanes and balanced winner reduction.
4. Machine-readable identities, hashes, cache behavior, and Pareto history supporting replay and claim auditing.

2. Adaptive- K Contract

For mode m , HM instrumentation forms the stable ranking

$$C_{\text{rough}}[m] = \text{SATD}[m] + \text{bits}[m]\sqrt{\lambda}. \quad (1)$$

Our Full-RDO anchor evaluates all 35 ranked modes. The frozen `adaptive-threshold.v1` normalizes the minimum rough cost by luma block area and sample range, then chooses

$$K = \begin{cases} 4, & \hat{C}_{\min} \leq 0.04, \\ 16, & 0.04 < \hat{C}_{\min} \leq 0.20, \\ 35, & \hat{C}_{\min} > 0.20. \end{cases} \quad (2)$$

The controller does not consume P , scheduler state, QP, or activity. This separation prevents hardware exploration from silently changing the selected software policy. Relative-SATD and hardware-aware variants receive distinct experiment versions rather than being folded into the frozen controller.

Selection used two deterministic synthetic workloads. Evaluation then ran without retuning on three predeclared held-out synthetic families. All matched curves use HM-16.20 commit `22178e370178133438c0339f57b3b3a29f112909`, Main 8-bit 4:2:0 All-Intra, four frames, and QP 22/27/32/37.

3. End-to-End Co-Design Architecture




Figure 1. Complete CHIA-RDO loop. Dominated and failed points remain in history, preventing repeated proposals and turning backend failures into the next architecture decision.

Native `ChiaFunction` nodes pass Ray object references asynchronously. The deterministic controller consumes completed records, excludes existing experiment IDs, and rewards novelty and Pareto-boundary exploration. The initial graph screened 72 (P , depth, buffer, width) points. A multi-wave campaign accumulated 114 hardware/history records, 107 successful logical experiments, 11 matched HM curves (132 encodes), and a 10-point Pareto set. The calibrated graph later screened 832 combinations over $P = 1 \dots 16$, depth 2/3, and cost width 31–56. It produced 16 Pareto points; deliberate replay returned 832/832 cache hits with no new evaluation.

4. RTL Architecture, Checked Against Source

The final top level latches an event containing $K \leq 35$ ranked mode descriptors, QP, Q16.16 λ , references, and a 4×4 original luma block. A ready/valid scheduler launches candidate $bP + l$ to lane l and masks the partial final batch. Up to 16 lanes operate concurrently; each lane is internally serial and accepts another candidate only after retirement.




Figure 2. Implemented RTL hierarchy and lane FSM. Exact HM candidate rate is external; hardware CABAC/context evolution is out of scope.

The lane implements planar, DC, and angular modes 0–34; 4×4 8-bit prediction/residual; two-pass HEVC integer DST; scalar quantization with RDOQ, transform skip, sign hiding, and scaling lists disabled; inverse transform and clipped reconstruction; SSE; and fixed-point RD cost. Depth 3 separates quantization from dequantization. One candidate takes 132 cycles and a batch retires in 134 cycles without backpressure.

The reduction compares valid pairs lexicographically: lower RD cost wins and an exact tie selects lower stable rank. In the P16 revision, a $\log_2 P$ balanced tree replaces the P-deep scan. The 31-bit cost is safe for the observed maximum 1,312,532,898 only; it is not an arbitrary-input proof.

Verification evidence

Final campaigns executed 121,600 vector/candidate checks. The depth-3 serial corpus alone contains 8,960 vectors covering all 35 modes and four QPs; the tree campaign adds 10,560 P1/P8/P16 evaluations including backpressure. Stage values, order, winners, cycle counts, dropped/duplicate checks, and interfaces all passed with zero failures.

5. Coding and Hardware Evaluation

Coding trade-off

Method / corpus	Candidate rule	Avg. K	RDO red.	Retention	BD-rate
Full-RDO anchor	All 35 rough-ranked modes	35.000	0.00%	100.00%	0.000%
Fixed- K 8 / discovery	First 8 ranked modes	8.000	77.14%	68.05%	+3.574%
Adaptive-v0 / discovery	Confidence gap; $K = 4/8/16$	11.568	66.95%	75.67%	+2.866%
Relative-SATD / discovery	Threshold count, quantized K	25.624	26.79%	96.07%	+0.528%
Frozen threshold / discovery	Normalized rough cost; $K = 4/16/35$	18.651	46.71%	90.25%	+0.755%
Frozen threshold / held-out	Same policy, no retuning	21.628	38.21%	97.91%	-0.026%
HW-aware Track-B / P8	Joint quality/cycle objective	16.456	52.98%	88.02%	-0.054%

Table 1. Project results relative to the Full-RDO anchor. The held-out set contains three synthetic families; negative aggregate BD-rate is not claimed as general quality improvement.

The frozen policy responds to content. It removes 13.95% and 13.30% of RDO work on mixed-texture and repetitive families, but 87.36% on smooth/edge content. Worst individual-run winner retention is 91.00%. On the held-out aggregate, Fixed- K 16 removes more evaluations (54.29%) but retains only 85.70% of Full-RDO winners. Track-B is a separately versioned, more aggressive hardware-oriented point; it is not evidence that the frozen policy was altered after selection.

The Gwon-Choi relative-SATD Bayesian method reports 31.54% average encoding time reduction and +0.93% BD-rate against HM-16.6 under its own Intra Main condition [2]. It starts from stock-HM RMD modes and includes a minimum-risk Bayesian classifier. Our relative-SATD baseline instead starts from all 35 modes and uses threshold counting only, so the published result is context rather than a direct comparison.

6. Physical Feedback as an Agent Input

The first P-way baseline replicated a monolithic candidate datapath. At P1, the critical path traversed prediction, residual, both forward-transform passes, quant/dequant, both inverse passes, reconstruction, and distortion: 170 logic levels, including 69 CARRY8 levels. P2 reached 175 levels, and P4 exceeded LUT/CARRY capacity before placement.

The next proposal serialized work within each lane and inserted state boundaries between major arithmetic phases. This raised Fmax and reduced per-lane area, allowing P16 to route, but a serial 16-way winner scan became an 80-level path. The final balanced tree reduced that path; the routed critical path is now 39 levels inside forward-transform accumulation and ends at a residual-sample RAM input.

Measured interpretation

For a failed 5 ns timing constraint, we report

$$f_{\text{derived}} = \frac{1}{5 \text{ ns} - \text{WNS}}, \quad T_{\text{kernel}} = \frac{P f_{\text{derived}}}{C_{\text{batch}}}. \quad (3)$$

This is a post-route comparison proxy. It is neither timing closure nor encoder frames/s. The selected P16 design routes with WNS = -9.680 ns, hence $f_{\text{derived}} = 68.120$ MHz and $T_{\text{kernel}} = 8.134$ M candidates/s for 134-cycle batches.

Resource budget

On the KV260 device (117,120 LUTs, 1,248 DSPs), the final design uses 52.26% of LUTs and 65.38% of DSPs. Against the project’s 70%-of-device integration budget, this is 74.66% of the LUT allowance and 93.47% of the DSP allowance, leaving 57 DSPs. P16 is the throughput endpoint of the Pareto frontier, not the minimum-area or minimum-power point.

7. Routed Microarchitecture Comparison

Architecture	<i>P</i>	Result	LUT	DSP	WNS	Derived Fmax	Kernel rate
Monolithic combinational	1	Routed	36,454	178	−60.395 ns	15.292 MHz	3.823 M/s
Monolithic combinational	2	Routed	70,817	356	−64.653 ns	14.357 MHz	7.178 M/s
Monolithic combinational	4	Placement failed	140,102 [†]	712	N/A	N/A	N/A
Shared serial, d3, w31	8	Routed	30,813	408	−9.620 ns	68.399 MHz	4.084 M/s
Shared serial, P16 scan	16	Routed	60,232	816	−19.092 ns	41.508 MHz	4.956 M/s
Shared serial, P16 tree	16	Routed	61,207	816	−9.680 ns	68.120 MHz	8.134 M/s

Table 2. Vivado 2024.2, xck26-sfvc784-2LV-c, 5 ns out-of-context constraint. [†]Synthesized LUT count; placement did not run. No failed value is interpolated.

The balanced tree costs 975 LUTs and 45 FFs over the P16 serial scan, leaves DSP count unchanged, and reduces estimated power from 2.338 W to 2.288 W. It improves derived Fmax by 64.1% and kernel rate by 64.1%. Relative to the monolithic P1 proxy, its candidate rate is 2.128×; relative to P2 it is 1.133×. The calibrated CHIA model predicts P16 within 0.14% for LUT, 0.45% for Fmax/throughput, and 0.43% for power.

8. Reproducibility and Deliverables

Local/GCP parity produced identical bitstream, reconstruction, and trace hashes for a reference encode. GCP workers later replayed RTL campaigns with zero failures. Post-run Compute Engine queries found no remaining instances, disks, or reserved addresses. Actual billed spend and promotional-credit balance remain unavailable and are not replaced by estimates.

The artifact preserves policy versions, manifests, source and corpus hashes, raw Vivado reports, machine-readable result schemas, and Pareto history. Principal deliverables are:

- CHIA proposal/evaluation graphs and deterministic policy controller;
- pinned HM patches, matched quality curves, and synthetic fixtures;
- parameterized SystemVerilog, golden vectors, and Verilator campaigns;
- analytical replay, model calibration, routed reports, and parsers.

9. Limitations and Next Step

The coding study uses small deterministic synthetic All-Intra workloads, not standard natural-video classes. RTL is limited to 4×4 8-bit luma and excludes chroma, inter

prediction, larger transforms, RDOQ, syntax coding, and hardware CABAC/context evolution. Exact HM candidate bit counts are supplied externally. The 31-bit cost width is corpus-qualified, not formally bounded for arbitrary input.

All physical runs are out of context and omit board-level clocks, pins, AXI, and encoder integration. Most importantly, the selected implementation has negative WNS and does not close 200 MHz. The next loop should pipeline or restructure forward-transform accumulation, now the measured critical path, then repeat regression and route before any board-frequency claim.

Within these limits, CHIA-RDO demonstrates the intended agentic co-design behavior: it preserves a software-quality contract, searches reproducibly, checks a bit-exact kernel, and uses failed backend evidence to produce a specific microarchitectural revision rather than merely requesting more parallelism.

References

- [1] CHIA, “An Open Framework for Agile and Principled Hardware/Software Co-Design Research,” arXiv:2606.27350, 2026.
- [2] D. Gwon and H. Choi, “Relative SATD-based Minimum Risk Bayesian Framework for Fast Intra Decision of HEVC,” *KSII Trans. Internet Inf. Syst.*, vol. 13, no. 1, pp. 385–405, Jan. 2019, doi:10.3837/tiis.2019.01.022.
- [3] ITU-T/ISO/IEC, “High Efficiency Video Coding Test Model HM-16.20.”
- [4] CHIA-RDO artifact, <https://github.com/duongnguyenhieu/CHIA-RDO>, release v1.0.0.

AI Assistance Acknowledgement

The authors used OpenAI OpenCode to assist with code review, artifact organization, and manuscript editing. The human authors reviewed the output and are responsible for the paper’s contents, tone, and quality.